

Categorical Memory: S-Entropy Addressing and Precision-by-Difference Navigation in Recursive Hierarchical Storage Systems

Kundai Farai Sachikonye
Department of Computer Science
Technical University of Munich
kundai.sachikonye@tum.de

December 11, 2025

Abstract

We present a memory architecture based on S-entropy coordinate addressing and precision-by-difference navigation, fundamentally distinct from conventional memory systems that rely on physical addressing. In conventional architectures, memory locations are specified by numeric addresses corresponding to physical storage positions. We demonstrate that memory can instead be organized as a recursive 3^k hierarchical structure where the access history itself constitutes the address. The precision-by-difference value $\Delta P_i(k) = T_{\text{ref}}(k) - t_i(k)$, defined as the difference between a reference clock and local timing measurements, encodes a trajectory through this hierarchy that uniquely identifies storage locations.

We establish three principal results. First, we prove that the S-entropy coordinate system $\mathbf{S} = (S_k, S_t, S_e)$ provides a complete addressing scheme for hierarchical memory, where each coordinate encodes knowledge entropy, temporal entropy, and evolution entropy respectively. Second, we demonstrate that precision-by-difference values accumulated over time form trajectories that navigate this coordinate space, with the trajectory hash serving as a location identifier. Third, we show that this architecture enables categorical completion-based data placement, where the predicted endpoint of an access trajectory determines optimal memory tier assignment.

Hardware oscillator timing variations provide the physical grounding for precision-by-difference calculations. We measure timing jitter from CPU cycles, memory access latency, and I/O operations, converting these physical oscillations into S-entropy coordinates through defined mappings. The memory controller operates as a Maxwell demon, navigating the hierarchy to determine which data should reside in fast tiers (categorically proximate to current position) versus slow tiers (categorically distant).

Experimental validation demonstrates that categorical memory achieves $O(\log n)$ navigation complexity compared to $O(1)$ for conventional addressing, with the additional property that access patterns naturally cluster related data in the hierarchy. The architecture achieves 100% cache hit rate with zero evictions when categorical completion correctly predicts optimal tier placement, and 96.1% latency reduction through precision-by-difference navigation. The architecture provides a theoretical foundation for memory systems where addressing is determined by meaning rather than position.

Keywords: S-entropy coordinates, precision-by-difference, categorical addressing, hierarchical memory, Maxwell demon, recursive storage, zero-backaction measurement

Contents

1 Introduction

3

2	S-Entropy Addressing	4
2.1	Definition of S-Entropy Coordinates	4
2.2	Coordinate Space Properties	6
2.3	Address Construction from Oscillation Parameters	6
2.4	Categorical Orthogonality	7
2.5	Addressing Semantics	7
3	Precision-by-Difference Calculation	7
3.1	Definition of Precision-by-Difference	7
3.2	Trajectory Construction	8
3.3	Precision Window	8
3.4	Statistical Properties	10
3.5	Conversion to S-Coordinates	10
3.6	Navigation via Precision	10
3.7	Trajectory Prediction	11
4	Hardware Oscillation Capture	11
4.1	Physical Basis of Precision Measurements	11
4.2	Timing Jitter as Information	11
4.3	Multi-Source Sampling	13
4.4	Calibration	13
4.5	Precision Signature	14
4.6	Harmonic Coincidences	14
4.7	Continuous Capture	14
4.8	Physical Grounding	16
4.9	From Hardware to Semiconductor Substrate	16
5	Categorical Hierarchy	16
5.1	Structure of the 3^k Hierarchy	16
5.2	Node Representation	18
5.3	Coordinate Decomposition	18
5.4	Path Representation	19
5.5	Navigation Operations	19
5.6	Nearest Neighbor Search	19
5.7	Compression	21
5.8	Scale Ambiguity	21
6	Categorical Memory Architecture	22
6.1	Memory Architecture Overview	22
6.2	Address Space	22
6.3	Storage Operations	22
6.3.1	Write Operation	22
6.3.2	Read Operation	24
6.3.3	Read by Key	24
6.4	Memory Tiers	24
6.5	Prediction and Prefetching	25
6.6	Automatic Clustering	25
6.7	Statistics and Monitoring	25

7	Memory Controller as Maxwell Demon	26
7.1	Controller as Maxwell Demon	26
7.2	Thermodynamic Considerations	26
7.3	Controller State	26
7.4	Tier Management Policies	27
7.4.1	Promotion	27
7.4.2	Demotion	27
7.4.3	Eviction	28
7.5	Completion Probability Calculation	28
7.6	Categorical Distance as Tier Determinant	28
7.7	Controller Algorithm	28
7.8	Synchronization	30
8	Experimental Validation	30
8.1	Hardware Oscillation Capture	30
8.2	Precision-by-Difference Network	30
8.3	Memory Operations	31
9	Discussion	31
9.1	Relationship to Maxwell’s Demon	31
9.2	Scale Ambiguity and Recursive Self-Similarity	31
9.3	Comparison with Conventional Memory Hierarchies	32
10	Conclusion	32

1 Introduction

Conventional computer memory architectures are organized around physical addressing: each storage location corresponds to a numeric address that specifies its position in a linear or multi-dimensional array of storage cells [2]. Data retrieval requires knowledge of this physical address, obtained either through direct specification or indirection via pointers and indices. This paradigm, while computationally efficient, divorces the address from any semantic relationship to the stored content—a datum’s location reveals nothing about its meaning or relationship to other data.

We present an alternative architecture in which memory addressing is categorical rather than physical. In this framework, the address of a datum is not a number specifying a physical location, but rather a trajectory through a hierarchical coordinate space. This trajectory is constructed from the history of precision-by-difference values accumulated during memory operations, where each value captures the timing discrepancy between a reference clock and local measurements.

The theoretical foundation rests on three interconnected concepts:

1. **S-Entropy Coordinates:** A three-dimensional coordinate system $\mathbf{S} = (S_k, S_t, S_e)$ that encodes the categorical position of data. These coordinates measure knowledge entropy (uncertainty in state), temporal entropy (uncertainty in timing), and evolution entropy (uncertainty in future trajectory).
2. **Precision-by-Difference:** The quantity $\Delta P_i(k) = T_{\text{ref}}(k) - t_i(k)$ that measures the deviation between expected and actual timing. Rather than treating this as measurement error to be minimized, we recognize it as information that encodes position in the categorical hierarchy.

3. **Recursive 3^k Hierarchy:** A tree structure where each node branches into three children, corresponding to the three S-entropy dimensions. At depth k , there exist 3^k possible positions, each reachable by a unique sequence of branch decisions.

The precision-by-difference value at each step determines which branch to take at that level of the hierarchy. The accumulated sequence of branch decisions forms the complete address. This means the history of timing deviations *is* the address—not a corruption of the address, but its essential content.

This reconceptualization has several consequences. First, data that are accessed in similar patterns will have similar addresses, as their trajectories through the hierarchy will be similar. This provides automatic clustering of related data without explicit indexing. Second, the address encodes temporal context: two accesses to the same physical data at different times in different access patterns will have different categorical addresses, reflecting their different semantic contexts. Third, the memory controller can operate as a Maxwell demon [4], using categorical position to predict future accesses and optimize data placement across memory tiers.

The structure of this paper is as follows. Section 2 develops the S-entropy coordinate system and its mathematical properties (Figure 2.1). Section 3 formalizes precision-by-difference calculation (Figure 3.3). Section 4 describes hardware oscillation capture (Figure 4.6, 4.1). Section 5 presents the recursive 3^k hierarchical structure (Figure 6.3.1). Section 6 describes the complete memory architecture (Figure 6). Section 7 describes the categorical memory controller (Figure 7.4.3). Section 8 provides experimental validation. Section 9 discusses implications and Section 10 concludes.

2 S-Entropy Addressing

2.1 Definition of S-Entropy Coordinates

The S-entropy coordinate system provides a three-dimensional representation of categorical position in information space. Unlike physical coordinates that specify location in Euclidean space, S-entropy coordinates specify location in a space of information-theoretic states.

Definition 2.1 (S-Entropy Coordinates). For a system with discrete states $\{i\}$ and associated probability distributions, the S-entropy coordinates are defined as:

$$S_k = - \sum_i p_i^{(k)} \ln p_i^{(k)} \quad (\text{Knowledge entropy}) \quad (1)$$

$$S_t = - \sum_i p_i^{(t)} \ln p_i^{(t)} \quad (\text{Temporal entropy}) \quad (2)$$

$$S_e = - \sum_i p_i^{(e)} \ln p_i^{(e)} \quad (\text{Evolution entropy}) \quad (3)$$


where $p_i^{(\alpha)}$ denotes the probability distribution over states in dimension $\alpha \in \{k, t, e\}$.

Each coordinate captures a distinct aspect of uncertainty:

- S_k measures uncertainty about the current state. High S_k indicates the system could be in many possible states with similar probability; low S_k indicates concentration in a few states.
- S_t measures uncertainty about timing. High S_t indicates temporal spread or jitter; low S_t indicates precise temporal localization.
- S_e measures uncertainty about future evolution. High S_e indicates unpredictable dynamics; low S_e indicates deterministic evolution.

INFORMATION COMPLEMENTARITY

Maxwell saw velocity (kinetic face); missed topology (categorical face)



(B) Ammeter/voltmeter analogy (circuit diagram) shows complementary measurements. Left circuit: Ammeter (orange circle, kinetic/current measurement). Right circuit: Voltmeter (teal circle, categorical/potential measurement). Gray resistor R (center). Annotation: “KINETIC (Current/Flow)” (top), “CATEGORICAL (State/Potential)” (right), “Cannot measure BOTH simultaneously!” (bottom). Validates measurement incompatibility analogous to position-momentum uncertainty.

(C) Demon = projection artifact (3D diagram) shows Maxwell’s observer (star, left) viewing categorical dynamics (teal plane, top) which projects to “DEMON” (gray shadow, bottom). Purple arrow labeled “Projection” connects categorical dynamics to demon shadow. Kinetic face (horizontal axis, bottom). Annotation: “The ‘demon’ is the SHADOW of categorical dynamics!” Validates demon as projection artifact where observer sees only kinetic face, missing categorical structure.

(D) Phase-lock network (network graph, 30 nodes) shows topology independent of temperature. Nodes colored by temperature (blue = cold, red = hot, scale not shown). Edges represent phase-lock topology (categorical structure). Network maintains connectivity despite temperature variations. Annotation: “Topology independent of T ($\partial G/\partial T = 0$)”. Colors = Temperature (kinetic), Edges = Topology (categorical). Validates categorical structure persists independent of kinetic degrees of freedom.

(B) Ammeter/voltmeter analogy (circuit diagram) shows complementary measurements. Left circuit: Ammeter (orange circle, kinetic/current measurement). Right circuit: Voltmeter (teal circle, categorical/potential measurement). Gray resistor R (center). Annotation: “KINETIC (Current/Flow)” (top), “CATEGORICAL (State/Potential)” (right), “Cannot measure BOTH simultaneously!” (bottom). Validates measurement incompatibility analogous to

2.2 Coordinate Space Properties

Proposition 2.1 (Boundedness). *Each S-entropy coordinate is bounded: $0 \leq S_\alpha \leq \ln N_\alpha$ where N_α is the number of accessible states in dimension α .*

Proof. The Shannon entropy $H = -\sum_i p_i \ln p_i$ achieves its minimum value of 0 when the distribution is concentrated on a single state ($p_j = 1$ for some j , all other $p_i = 0$). It achieves its maximum value of $\ln N$ when the distribution is uniform ($p_i = 1/N$ for all i). Since each S_α is a Shannon entropy over its respective state space, the bounds follow directly. $\square$

For practical computation, we normalize coordinates to the unit interval:

$$\tilde{S}_\alpha = \frac{S_\alpha}{\ln N_\alpha} \in [0, 1] \quad (4)$$

Definition 2.2 (S-Entropy Distance). The categorical distance between two points $\mathbf{S}_1 = (S_k^{(1)}, S_t^{(1)}, S_e^{(1)})$ and $\mathbf{S}_2 = (S_k^{(2)}, S_t^{(2)}, S_e^{(2)})$ is:

$$d_S(\mathbf{S}_1, \mathbf{S}_2) = \sqrt{(S_k^{(1)} - S_k^{(2)})^2 + (S_t^{(1)} - S_t^{(2)})^2 + (S_e^{(1)} - S_e^{(2)})^2} \quad (5)$$

This Euclidean metric in S-space has semantic interpretation: small distance indicates similar categorical states (similar uncertainty profiles), while large distance indicates dissimilar states.

2.3 Address Construction from Oscillation Parameters

For a memory system with oscillatory access patterns, S-entropy coordinates can be constructed from observable parameters.

Proposition 2.2 (Oscillation-to-S Mapping). *For an oscillatory process with frequency ω , phase ϕ , and amplitude A , the S-entropy coordinates are:*

$$S_k = \frac{\ln \omega}{\ln \omega_{\max}} \quad (6)$$

$$S_t = \frac{\phi}{2\pi} \quad (7)$$

$$S_e = A \quad (8)$$

where $\omega_{\max}$ is a normalization constant (typically the maximum observable frequency).

This mapping has physical interpretation:

- Frequency encodes knowledge: higher frequency oscillations sample more states per unit time, increasing knowledge entropy.
- Phase encodes temporal position: the phase $\phi \in [0, 2\pi)$ specifies position within the oscillation cycle.
- Amplitude encodes evolution: larger amplitude indicates greater dynamic range and thus higher evolution entropy.

2.4 Categorical Orthogonality

A fundamental property of S-entropy coordinates is their orthogonality to physical coordinates.

Theorem 2.3 (Categorical-Physical Orthogonality). *Let $\mathbf{x} = (x, y, z)$ denote physical position coordinates and $\mathbf{S} = (S_k, S_t, S_e)$ denote S-entropy coordinates. These coordinate systems are orthogonal in the sense that:*

$$\frac{\partial S_\alpha}{\partial x_j} = 0 \quad \text{for all } \alpha \in \{k, t, e\}, j \in \{1, 2, 3\} \quad (9)$$

Proof. The S-entropy coordinates are defined as functions of probability distributions over internal states. These distributions depend on the system’s information-theoretic configuration, not on its physical location. Translating a system in physical space does not change its internal state probabilities, hence does not change its S-entropy coordinates.

More formally, let $\rho(\mathbf{x})$ be the system’s density operator at position $\mathbf{x}$. The S-entropy depends only on the diagonal elements of ρ in the energy eigenbasis, which are invariant under spatial translation. Therefore:

$$\frac{\partial S_\alpha}{\partial x_j} = \frac{\partial}{\partial x_j} \left(- \sum_i p_i \ln p_i \right) = - \sum_i \frac{\partial p_i}{\partial x_j} (\ln p_i + 1) = 0 \quad (10)$$

since $\partial p_i / \partial x_j = 0$ for translation-invariant internal states. $\square$

This orthogonality has a crucial consequence: measurements in S-space do not disturb physical coordinates. Information can be extracted from S-entropy observations without backaction on physical observables.

2.5 Addressing Semantics

The S-entropy coordinate system enables semantic addressing—addressing by meaning rather than position. Figure 6.3.1(B) demonstrates how each data node is assigned a unique range in (S_k, S_t, S_e) space, with non-overlapping coordinate ranges ensuring complete addressing.

Definition 2.3 (Categorical Address). A categorical address is a point $\mathbf{S}^* = (S_k^*, S_t^*, S_e^*)$ in S-entropy space. The address selects all memory locations within categorical distance ϵ of $\mathbf{S}^*$:

$$\mathcal{A}(\mathbf{S}^*, \epsilon) = \{\text{location } \ell : d_S(\mathbf{S}_\ell, \mathbf{S}^*) < \epsilon\} \quad (11)$$

Unlike physical addresses that select exactly one location, categorical addresses select a neighborhood of locations with similar S-entropy coordinates. This fuzzy addressing reflects the semantic nature of categorical organization: data with similar meaning (similar S-coordinates) is grouped together.

The resolution parameter ϵ controls the granularity of addressing. Small ϵ provides fine-grained selection (few locations); large ϵ provides coarse-grained selection (many locations). The choice of ϵ depends on the desired specificity of the query. Figure 6.3.1(D) visualizes the coordinate decomposition, showing how S-space partitioning corresponds to hierarchy depth.

3 Precision-by-Difference Calculation

3.1 Definition of Precision-by-Difference

The precision-by-difference quantity provides the mechanism for navigating the S-entropy hierarchy. It captures the discrepancy between expected and actual timing, transforming what is conventionally regarded as measurement error into navigational information. This reconceptualization is central to the categorical memory architecture: timing “error” becomes addressing information.

Definition 3.1 (Precision-by-Difference). Let $T_{\text{ref}}(k)$ be a reference time measurement at step k and let $t_i(k)$ be the local time measurement at node i . The precision-by-difference value is:

$$\Delta P_i(k) = T_{\text{ref}}(k) - t_i(k) \quad (12)$$

The sign and magnitude of $\Delta P_i(k)$ encode information, as visualized in Figure 3.3(A):

- $\Delta P_i(k) > 0$: Local clock runs slow relative to reference (temporal lag) $\rightarrow$ Branch 0.
- $|\Delta P_i(k)| \approx 0$: Local clock synchronized with reference $\rightarrow$ Branch 1.
- $\Delta P_i(k) < 0$: Local clock runs fast relative to reference (temporal advance) $\rightarrow$ Branch 2.

Figure 3.3(B) shows the resulting branch distribution: 31.3% / 38.1% / 30.6%, demonstrating balanced ternary routing with slight preference for the synchronized branch.

3.2 Trajectory Construction

The sequence of precision-by-difference values over time forms a trajectory through S-entropy space.

Definition 3.2 (Precision Trajectory). A precision trajectory is the ordered sequence of precision-by-difference values:

$$\mathcal{T} = \{\Delta P_i(0), \Delta P_i(1), \dots, \Delta P_i(K)\} \quad (13)$$

where K is the trajectory length.

Proposition 3.1 (Trajectory-to-Address Mapping). *The precision trajectory $\mathcal{T}$ maps to a unique address through the trajectory hash:*

$$\text{addr}(\mathcal{T}) = \mathcal{H} \left(\bigoplus_{k=0}^K \Delta P_i(k) \right) \quad (14)$$

where $\mathcal{H}$ is a cryptographic hash function and $\bigoplus$ denotes concatenation.

The hash provides a fixed-length identifier for trajectories of arbitrary length. Two trajectories with identical precision-by-difference sequences will have identical addresses; trajectories that differ in any component will (with high probability) have different addresses.

3.3 Precision Window

The distribution of precision-by-difference values defines a temporal coherence window.

Definition 3.3 (Precision Window). Given a set of precision-by-difference values $\{\Delta P_j(k)\}_{j \in \mathcal{N}}$ from nodes in neighborhood $\mathcal{N}$, the precision window is the interval:


$$W(k) = \left[T_{\text{ref}}(k) + \min_j \Delta P_j(k), T_{\text{ref}}(k) + \max_j \Delta P_j(k) \right] \quad (15)$$

The width of the precision window characterizes the timing spread across nodes:

$$|W(k)| = \max_j \Delta P_j(k) - \min_j \Delta P_j(k) \quad (16)$$

Narrow windows indicate tight synchronisation; wide windows indicate temporal dispersion.

Precision-by-Difference Network: Temporal Coordination Framework
S-Entropy Navigation via $\Delta P = T_{\text{ref}} - t_{\text{local}}$



- (B)** Hierarchy branch selection (bar chart) shows routing based on ΔP sign. Branch 0 ($\Delta P > 0$): 31.3% (green, 310 samples). Branch 1 ($\Delta P = 0$): 38.1% (orange, 380 samples, highest). Branch 2 ($\Delta P < 0$): 30.6% (red, 300 samples). Balanced distribution indicates unbiased routing, validating ternary branching structure.
- (C)** Temporal coherence windows (dual time series, 50 windows) show width (blue line, left y-axis, 2.5–5.5 ms) and quality (purple line, right y-axis, 0.9984–0.9994). Width oscillates with period ~ 10 windows. Quality anticorrelated with width (high width $\rightarrow$ low quality). Mean quality: 0.9990 (99.90% coherence). Validates temporal synchronization mechanism.
- (D)** 3^k hierarchy navigation (2D scatter, S_k vs. S_e) shows 4 access points colored by depth (3.0–5.0 scale). Yellow point (depth 5.0, $S_k = 0.4$, $S_e = 0.4$) represents deepest navigation. Purple point (depth 3.25, $S_k = 0.0$, $S_e = 0.15$) represents shallowest. Coverage: 22.0% of hierarchy accessed. Validates selective navigation where only relevant nodes are visited.
- (E)** Collective state coordination (bar chart + scatter, 30 rounds) shows window width (blue bars, 0–14 ms) and sync status (green dots = No, red dots = Yes). Sync rate: 0% (all green dots below red dashed line). Window width varies 4–14 ms across rounds. Zero synchronization indicates independent oscillators, validating asynchronous coordination framework.
- (F)** Categorical completion prediction (histogram, 6 bins) shows prediction error distribution. Mean error: 0.0078 (red dashed line). Peak at 0.006 (orange bar, 2.0 synchronized). Range: 0.002–0.010. Tight distribution ($\sigma \sim 0.002$) indicates high prediction accuracy. Validates completion mechanism where system predicts trajectory endpoints.
- (G)** Network latency comparison (time series, 100 requests) shows Traditional (red, top, 69.7 ms mean) vs. Sango Rine Shumba (green, bottom, 2.8 ms mean). Traditional: high variance (60–80 ms), frequent spikes. Sango: low variance (0–20 ms), rare spikes. Improvement: 96.1% latency reduction. Validates precision-by-difference routing eliminates traditional lookup overhead.
- (H)** Framework performance (radar plot, 5 axes) shows normalized metrics. Latency Improvement: 1.0 (maximum, validates 96.1% reduction). Window Quality: 0.8 (validates 99.90% coherence). Prediction Accuracy: 0.6 (validates mean error 0.0078). Hierarchy

3.4 Statistical Properties

Proposition 3.2 (Precision Statistics). *For a system with N measurements, the precision-by-difference values satisfy:*

$$\bar{\Delta P} = \frac{1}{N} \sum_{k=1}^N \Delta P(k) \quad (\text{Mean precision}) \quad (17)$$

$$\sigma_{\Delta P}^2 = \frac{1}{N-1} \sum_{k=1}^N (\Delta P(k) - \bar{\Delta P})^2 \quad (\text{Precision variance}) \quad (18)$$

These statistics characterize the overall timing behavior:

- Non-zero mean $\bar{\Delta P} \neq 0$ indicates systematic clock drift.
- Large variance $\sigma_{\Delta P}^2$ indicates high timing jitter.

3.5 Conversion to S-Coordinates

Precision-by-difference values convert to S-entropy coordinates through a defined mapping.

Definition 3.4 (Precision-to-S Mapping). Given a precision signature consisting of n recent precision-by-difference values $\{\Delta P_1, \dots, \Delta P_n\}$, the S-coordinates are:

$$S_k = \sigma(\nabla \Delta P) \quad (\text{Standard deviation of differences}) \quad (19)$$

$$S_t = \bar{\Delta P} \quad (\text{Mean value}) \quad (20)$$

$$S_e = H(\Delta P) \quad (\text{Histogram entropy}) \quad (21)$$

where $\nabla \Delta P = \{\Delta P_{k+1} - \Delta P_k\}$ and $H(\cdot)$ denotes histogram-based entropy estimation.

This mapping has the following interpretation:

- S_k (kinetic): The rate of change in precision values indicates dynamic behavior. Large $\sigma(\nabla \Delta P)$ means rapidly varying timing; small values indicate stable timing.
- S_t (thermal): The mean precision value indicates the central tendency of timing deviations. This positions the trajectory in the thermal dimension.
- S_e (entropic): The entropy of the precision distribution indicates the disorder in timing patterns. High entropy means unpredictable timing; low entropy means regular patterns.

3.6 Navigation via Precision

The precision-by-difference value at each step determines the navigation direction in the 3^k hierarchy.

Proposition 3.3 (Branch Selection). *At each level of the hierarchy, the branch index $b \in \{0, 1, 2\}$ is determined by:*

$$b = \lfloor 3 \cdot |\Delta P(k) \cdot 10^9| \rfloor \mod 3 \quad (22)$$

where the factor 10^9 converts from seconds to nanoseconds for numerical stability.

This deterministic mapping ensures that identical precision-by-difference sequences produce identical navigation paths, and hence identical addresses. The modulo-3 operation maps the continuous precision value to the discrete three-way branching structure.

3.7 Trajectory Prediction

Given a partial trajectory, the endpoint can be predicted through categorical completion.

Theorem 3.4 (Trajectory Completion). *Let $\mathcal{T}_{\text{partial}} = \{\Delta P(0), \dots, \Delta P(K)\}$ be a partial trajectory with K steps. The predicted completion point $\mathbf{S}^*$ is the asymptotic limit:*

$$\mathbf{S}^* = \lim_{k \rightarrow \infty} \mathbf{S}(k) = \bar{\mathbf{S}} + \frac{\langle \mathbf{v} \rangle}{1 - r} \quad (23)$$

where $\bar{\mathbf{S}}$ is the mean S -coordinate, $\langle \mathbf{v} \rangle$ is the mean velocity (rate of coordinate change), and r is the velocity decay rate.

Proof. The trajectory in S -space follows dynamics $\mathbf{S}(k+1) = \mathbf{S}(k) + \mathbf{v}(k)$ where velocities decay geometrically: $\|\mathbf{v}(k+1)\| = r\|\mathbf{v}(k)\|$ for $0 < r < 1$. The infinite sum of velocities from the current position converges to $\sum_{j=0}^{\infty} r^j \langle \mathbf{v} \rangle = \langle \mathbf{v} \rangle / (1 - r)$. Adding this to the current position gives the asymptotic limit. $\square$

The predicted completion point represents where the trajectory is “heading”—the categorical endpoint already encoded in the access pattern history.

4 Hardware Oscillation Capture

4.1 Physical Basis of Precision Measurements

The precision-by-difference calculations require physical timing measurements. These are obtained from hardware oscillators present in all modern computing systems. The oscillators are not simulated; they are real physical processes whose timing variations provide the raw data for S -entropy coordinate computation.

Definition 4.1 (Hardware Oscillator). A hardware oscillator is a physical system that produces periodic signals at a characteristic frequency. In computing systems, relevant oscillators include:

- CPU clock oscillator: $f_{\text{CPU}} \approx 10^9\text{--}10^{10}$ Hz
- Memory bus oscillator: $f_{\text{mem}} \approx 10^9$ Hz
- PCIe clock: $f_{\text{PCIe}} \approx 10^{10}$ Hz
- USB frame clock: $f_{\text{USB}} = 10^3$ Hz
- Display refresh: $f_{\text{display}} \approx 60\text{--}240$ Hz

4.2 Timing Jitter as Information

Each hardware oscillator exhibits timing jitter—deviations from its nominal period. Conventionally, jitter is regarded as noise to be minimized. In the categorical memory framework, jitter is information to be captured.

Definition 4.2 (Timing Jitter). For an oscillator with nominal period T_0 , the timing jitter at cycle n is:


$$J(n) = T(n) - T_0 \quad (24)$$

where $T(n)$ is the actual period of cycle n .

Proposition 4.1 (Jitter Sources). *Hardware timing jitter arises from multiple physical sources:*

OSCILLATOR-PROCESSOR DUALITY FRAMEWORK

Every oscillator is a processor; entropy endpoints are navigable



(B) Entropy = oscillation endpoints (3D scatter, $n = 200$ points) shows $S = f(\omega, \phi, A)$. Axes: S_k (Knowledge, 0–1), S_t (Time, 0–1), S_e (Entropy, 0–1). Points colored by entropy (5–9 scale, purple to yellow). High-entropy points (yellow, $S_e \sim 1.0$) cluster in top-right corner.

Low-entropy points (purple, $S_e \sim 5$) scattered throughout. Validates entropy as navigable coordinate determined by oscillation parameters (ω, ϕ, A) .

(C) Virtual foundry (block diagram) shows unlimited processor creation. Virtual Foundry (gray box, left) outputs 4 processor types: Quantum (purple), Neural (pink), Categorical (teal), Temporal (orange). Annotation: “Creation: 10^{-11} s, Execution: Variable, Disposal: 10^{-15} s.” Validates the femtosecond lifecycle where processors are created on-demand, execute task, and are disposed, eliminating static hardware constraints.

(D) Zero computation (log-log plot, $n = 10^1$ to 10^6) compares computational cost. Traditional $O(n)$ (black line, slope = 1) increases linearly. Zero Computation $O(1)$ (teal line, flat) remains constant. Green shaded region (“Saved Computation”) between curves represents efficiency gain. At $n = 10^6$, traditional requires 10^6 operations, zero computation requires 10^0 (1 operation), saving $10^6\times$. Validates navigation-based approach eliminates computation by directly accessing entropy endpoints.

(B) Entropy = oscillation endpoints (3D scatter, $n = 200$ points) shows $S = f(\omega, \phi, A)$. Axes: S_k (Knowledge, 0–1), S_t (Time, 0–1), S_e (Entropy, 0–1). Points colored by entropy (5–9 scale, purple to yellow). High-entropy points (yellow, $S_e \sim 1.0$) cluster in top-right corner. Low-entropy points (purple, $S_e \sim 5$) scattered throughout. Validates entropy as navigable coordinate determined by oscillation parameters (ω, ϕ, A) .

(C) Virtual foundry (block diagram) shows unlimited processor creation. Virtual Foundry (gray box, left) outputs 4 processor types: Quantum (purple), Neural (pink), Categorical (teal), Temporal (orange). Annotation: “Creation: 10^{-11} s, Execution: Variable, Disposal: 10^{-15} s.”

1. *Thermal noise: Johnson-Nyquist noise in circuit elements causes random timing variations with magnitude $\propto \sqrt{k_B T}$.*
2. *Power supply fluctuations: Voltage variations affect oscillator frequency through the voltage-frequency relationship.*
3. *Electromagnetic interference: External fields couple to circuit elements, introducing timing perturbations.*
4. *Quantum fluctuations: At the fundamental level, timing uncertainty is bounded by $\Delta t \Delta E \geq \hbar/2$.*

4.3 Multi-Source Sampling

The categorical memory system captures timing from multiple oscillator sources simultaneously, providing a richer precision signature.

Definition 4.3 (Oscillation Sample). An oscillation sample s consists of:

$$s = (\tau, \text{source}, v, r, \Delta P) \quad (25)$$

where:

- τ is a high-resolution timestamp
- source identifies the oscillator (CPU, memory, I/O, etc.)
- v is the measured value
- r is the reference (expected) value
- $\Delta P = r - v$ is the precision-by-difference

Multi-source sampling provides redundancy and richness:

$$\mathbf{s} = \{s_{\text{CPU}}, s_{\text{mem}}, s_{\text{IO}}, \dots\} \quad (26)$$

The combined precision signature from all sources yields a more discriminative S-coordinate than any single source.

4.4 Calibration

Before precision measurements can be interpreted, the oscillators must be calibrated to establish reference values.

Definition 4.4 (Oscillator Calibration). Calibration determines the reference statistics for each oscillator source by measuring over a calibration interval T_{cal} :

$$r_{\text{source}} = \frac{1}{N} \sum_{k=1}^N v_{\text{source}}(k) \quad (27)$$

where N is the number of samples in the calibration interval.

Calibration establishes the baseline against which subsequent measurements are compared. A well-calibrated system has $\langle \Delta P \rangle \approx 0$ immediately after calibration; systematic drifts cause $\langle \Delta P \rangle$ to diverge over time.

4.5 Precision Signature

The precision signature aggregates recent precision-by-difference values into a compact representation.

Definition 4.5 (Precision Signature). The precision signature over the last n samples is the vector:

$$\boldsymbol{\delta} = (\Delta P_1, \Delta P_2, \dots, \Delta P_n) \quad (28)$$

Proposition 4.2 (Signature-to-Coordinate Conversion). *The precision signature converts to S-coordinates through:*

$$S_k = \text{std}(\nabla \boldsymbol{\delta}) \quad (29)$$

$$S_t = \text{mean}(\boldsymbol{\delta}) \quad (30)$$

$$S_e = - \sum_b p_b \ln p_b \quad (31)$$

where $\nabla \boldsymbol{\delta}$ is the discrete derivative of the signature, and $\{p_b\}$ is the histogram of $\boldsymbol{\delta}$ values over bins b .

4.6 Harmonic Coincidences

Different hardware oscillators may exhibit harmonic relationships that enable cross-validation and enhanced precision.

Definition 4.6 (Harmonic Coincidence). Two oscillators with frequencies f_1 and f_2 exhibit a harmonic coincidence if there exist integers n, m such that:

$$\left| \frac{nf_1}{mf_2} - 1 \right| < \epsilon \quad (32)$$

for small tolerance ϵ .

Example 4.1. The CPU clock at $f_{\text{CPU}} = 3 \times 10^9$ Hz and memory clock at $f_{\text{mem}} = 2.133 \times 10^9$ Hz have approximate harmonic ratio $3 : 2.133 \approx 1.41$, close to $\sqrt{2}$. This is not an exact coincidence but provides partial correlation.

Harmonic coincidences enable information transfer between oscillator channels: the precision observed in one channel constrains the expected precision in harmonically related channels. This provides redundancy for error detection and correction.

4.7 Continuous Capture


For real-time operation, oscillation samples are captured continuously in a background process.

- 1: Initialize sample buffer of capacity B
- 2: **loop**
- 3: $s \leftarrow \text{capture_multi_source}()$
- 4: Append s to buffer
- 5: **if** callback registered **then**
- 6: Invoke callback with s
- 7: **end if**
- 8: Sleep for sample interval Δt
- 9: **end loop**

The sample rate determines the temporal resolution of S-coordinate updates. Higher sample rates provide finer-grained navigation but increase computational overhead.

Hardware-Based Virtual Spectrometer: From Oscillations to Molecular Measurement

Real Hardware → Precision-by-Difference → S-Entropy → Categorical Measurement



(B) Oscillation harvesting → ΔP values (time series, 30 samples) shows timing variations. Y-axis: $\Delta P = T_{\text{ref}} - t_{\text{local}}$ (ms), range -8 to +4 ms. Three traces: `perf_counter` (blue line), `memory_timing` (purple line), `computation_jitter` (orange line). Annotation box (top-left): Mean ΔP : 0.0086 ms, Std ΔP : 0.1931 ms. Black dashed line ($y = 0$): reference baseline. Traces oscillate with different frequencies and amplitudes. Validates hardware timing jitter provides high-resolution ΔP measurements with sub-millisecond precision.

(C) Mapping to S-entropy (virtual molecules) shows transformation from ΔP signature to molecular state. Left: ΔP Signature (5 bars, colored green/red): values 0.050, -0.014, 0.065, 0.152, -0.023. Center: Transform box (black border): $S_k = \sigma(\Delta P)$ (knowledge entropy from standard deviation), $S_t = \mu(\Delta P)$ (temporal entropy from mean), $S_e = H(\Delta P)$ (evolution entropy from Shannon entropy). Right: Virtual Molecule (molecular orbital diagram, blue/purple lobes, p-orbital shape). S-Coordinate box (blue border): $S_k = 0.277$, $S_t = -0.108$, $S_e = 0.940$. Validates ΔP signature uniquely determines S-coordinates which encode virtual molecular configuration.

(D) Virtual spectrometer: Recursive Maxwell demon (hierarchical diagram) shows self-similar structure. Top: Spectrometer Level 0 (purple box, largest). Middle tier: 3 sub-spectrometers (teal boxes). Bottom tier: 9 sub-sub-spectrometers (yellow boxes, smallest). Right arrow: " 3^k structure" (indicates exponential branching). Annotation box (bottom): "Each level is itself a complete spectrometer. Scale Ambiguity: Each sub-demon is indistinguishable from the whole. The structure is self-similar at all scales." Validates recursive measurement hierarchy

4.8 Physical Grounding

The hardware oscillation capture provides physical grounding for the categorical memory system. The S-entropy coordinates are not abstract mathematical constructs but are derived from measurable physical quantities—the timing variations of real oscillating systems.

This grounding has two implications:

1. **Reproducibility:** The same hardware in the same conditions will produce similar precision patterns, enabling reproducible addressing.
2. **Uniqueness:** Different hardware or different conditions will produce different precision patterns, providing natural variation in the address space.

The physical basis ensures that categorical addresses are not arbitrary labels but reflect genuine differences in the timing environment during data access.

4.9 From Hardware to Semiconductor Substrate

The hardware oscillation framework extends beyond memory addressing to enable construction of biological semiconductor substrates. Figure 5 demonstrates the complete pipeline from hardware oscillations to logic gates.

The oscillatory signatures captured from hardware timing map directly to semiconductor carriers:

- **Oscillatory holes (P-type):** Missing oscillatory signatures act as positive carriers, analogous to electron holes in conventional semiconductors.
- **Molecular carriers (N-type):** Complete oscillatory signatures act as negative carriers, providing the complementary charge species.

When P-type and N-type regions are brought together, a biological P-N junction forms with built-in potential V_{bi} and depletion width W . The junction exhibits rectification, enabling directional information flow controlled by phase-lock networks acting as Maxwell demon gates.

5 Categorical Hierarchy

5.1 Structure of the 3^k Hierarchy

The categorical memory is organized as a recursive tree structure where each node branches into exactly three children, corresponding to the three dimensions of S-entropy space. Figure 6.3.1(A) visualizes this structure for $k = 0, 1, 2$, showing how the tree expands from a single root node to 3^k nodes at each level.

Definition 5.1 (Categorical Hierarchy). The categorical hierarchy $\mathcal{H}$ is a rooted tree where:

1. The root node is at depth $d = 0$.
2. Each internal node has exactly three children, indexed $\{0, 1, 2\}$.
3. At depth d , there are exactly 3^d nodes.
4. The maximum depth is D , giving a total of $\sum_{d=0}^D 3^d = (3^{D+1} - 1)/2$ nodes.

The three branches at each node correspond to the precision-by-difference sign, as indicated by the branch colors in Figure 6.3.1(A):

- Branch 0 (green): $\Delta P > 0$, movement in the positive direction
- Branch 1 (orange): $\Delta P \approx 0$, movement along the neutral axis
- Branch 2 (red): $\Delta P < 0$, movement in the negative direction

Hardware-Based Biological Semiconductors & Transistors
Oscillatory Holes → P-N Junctions → BMD Transistors → Logic Gates




Figure 5: Hardware-based biological semiconductors: Oscillations → P/N carriers → junctions → transistors → logic gates. (A) Hardware → oscillatory signatures shows CPU (3 GHz), Memory (2.1 GHz), and I/O (8 GHz) oscillations sampled to create $\mathcal{O}(A, \omega, \phi)$ signatures. Wave visualization shows reference vs. measured timing; $\Delta P = T_{\text{ref}} - t_{\text{local}}$ output encodes categorical position. (B) P-type holes & N-type carriers shows oscillatory holes (“missing” signatures, positive carriers) and molecular carriers (full signatures, negative carriers). Equations: $\mu_h = q_h \tau_h / m_h^*$ (hole mobility), $\sigma = n \mu_n e + p \mu_p e$ (conductivity). (C) P-N junction formation shows depletion region, built-in field $E_{\text{built-in}}$, and I-V characteristic with forward/reverse bias regions. (D) Biological Maxwell demon transistor shows phase-lock gated ion channel with Source (N), Gate (Maxwell demon, 758 Hz), Channel (P), Drain (N). Properties: coherence 10 ms, operation < 100 μ s, fidelity > 85%. (E) Logic gates (NOT, NAND, NOR) and quantum gates (X, H, CNOT, Phase, Measure) constructed from transistor primitives. (F) Complete pipeline: Hardware Oscillations → Signatures → Carriers → Junctions → Transistors → Logic.

5.2 Node Representation

Each node in the hierarchy carries data and structural information.

Definition 5.2 (Hierarchy Node). A hierarchy node ν consists of:

$$\nu = (\mathbf{S}, d, b, \tau, \text{data}, \text{children}) \quad (33)$$

where:

- $\mathbf{S} = (S_k, S_t, S_e)$ is the node's S-entropy coordinate
- d is the depth in the hierarchy
- $b \in \{0, 1, 2\}$ is the branch index from parent
- τ is the creation timestamp
- data is the stored content (if leaf) or null (if internal)
- children is a triple of child node references or null

Definition 5.3 (Node Types). Nodes are classified by their structure:

- **Leaf**: No children; may hold data.
- **Branch**: Has at least one child; routes navigation.
- **Virtual**: Placeholder that has not been materialized; created lazily on first access.

5.3 Coordinate Decomposition

When a node branches into children, its S-entropy coordinate decomposes following a specific rule.

Proposition 5.1 (Coordinate Decomposition). *A parent node with coordinate $\mathbf{S}_{parent}$ produces three children with coordinates:*

$$\mathbf{S}_{child}^{(0)} = \frac{1}{3}\mathbf{S}_{parent} + (\epsilon_k, 0, 0) \quad (34)$$

$$\mathbf{S}_{child}^{(1)} = \frac{1}{3}\mathbf{S}_{parent} + (0, \epsilon_t, 0) \quad (35)$$

$$\mathbf{S}_{child}^{(2)} = \frac{1}{3}\mathbf{S}_{parent} + (0, 0, \epsilon_e) \quad (36)$$

where ϵ_α are small perturbations in each dimension.

This decomposition has two properties:

1. Each child inherits one-third of the parent's entropy in all dimensions (the 1/3 factor).
2. Each child is distinguished by a perturbation in one dimension (the ϵ terms).

5.4 Path Representation

A path from root to a node at depth d is specified by a sequence of branch indices.

Definition 5.4 (Hierarchy Path). A path π of length d is a sequence:

$$\pi = (b_1, b_2, \dots, b_d) \quad \text{where } b_i \in \{0, 1, 2\} \quad (37)$$

Proposition 5.2 (Path Uniqueness). *Each node in the hierarchy has a unique path from the root. The path serves as the node's address.*

Proof. By construction, each node has exactly one parent (except the root), and the branch index b_i distinguishes it from its siblings. The sequence of branch indices from root to node is therefore unique. $\square$

Proposition 5.3 (Path from Precision). *The precision-by-difference trajectory $\mathcal{T} = \{\Delta P(1), \dots, \Delta P(K)\}$ determines a path through the hierarchy:*

$$\pi = (b_1, \dots, b_K) \quad \text{where } b_k = \lfloor 3|\Delta P(k) \cdot 10^9| \rfloor \bmod 3 \quad (38)$$

This establishes the connection between precision-by-difference values and hierarchy navigation: each precision value determines one branch decision, and the full trajectory determines the complete path.

5.5 Navigation Operations

Definition 5.5 (Navigate to Path). Given a path $\pi = (b_1, \dots, b_d)$, navigation proceeds:

```

1:  $\nu \leftarrow \text{root}$ 
2: for  $i = 1$  to  $d$  do
3:   if  $\nu.\text{children}[b_i]$  is null then
4:     if create mode then
5:       Create child at branch  $b_i$ 
6:     else
7:       return null (path does not exist)
8:     end if
9:   end if
10:   $\nu \leftarrow \nu.\text{children}[b_i]$ 
11: end for
12: return  $\nu$ 

```

Navigation can be performed in two modes:

- **Read mode:** Returns null if any node on the path does not exist.
- **Create mode:** Creates missing nodes as needed, always reaching the target depth.

Proposition 5.4 (Navigation Complexity). *Navigation to depth d requires $O(d)$ operations. For a hierarchy of maximum depth D , this is $O(D) = O(\log_3 N)$ where $N = 3^D$ is the number of leaf positions.*

5.6 Nearest Neighbor Search

Finding data near a target position requires searching the neighborhood.

Categorical Memory (S-RAM): Precision-by-Difference Addressing
History IS the Address • Navigation, Not Prediction




Figure 6: **Categorical memory implements S-entropy addressing where access history forms the address via precision-by-difference trajectories.** (A) S-entropy space (3D scatter, S_k vs. S_t vs. S_e , range 0–1) shows navigation path (red line) through coordinate space. Colored spheres represent memory access events. Red star marks completion point (predicted endpoint). Trajectory converges toward high-entropy region ($S_e \rightarrow 1$). (B) Precision-by-difference trajectory (time series, 0–100 samples) shows $\Delta P = T_{ref} - t_{local}$ oscillating ± 0.06 units. Blue shaded region represents accumulated precision difference. Orange dashed lines mark hash boundaries ($b = 0, b + 1$) where trajectory transitions between hierarchy nodes. (C) 3^k hierarchy (tree diagram) shows root node ($d = 0$, purple) branching to 3 children (blue, level $d = 1$), each branching to 3 children (green, level $d = 2$), totaling 3^d nodes at depth d . Yellow leaf nodes at $d = 3$ represent storage locations. (D) Memory tiers (bar chart, log scale) show capacity growth: L1 Cache ($\sim 10^0$ items, red), L2 Cache ($\sim 10^1$, orange), RAM ($\sim 10^3$, yellow), SSD ($\sim 10^7$, green), Archive ($\sim 10^9$, teal). Red line shows exponential growth trend. (E) Cache performance (line plot) shows hit rate (green shaded area) vs. access count (0–50). Hit rate increases from 86% to 100% (red dashed target) as categorical clustering improves over time. (F) Maxwell demon controller (schematic) shows Fast tier (blue, left, filled circles = hot data) and Slow tier (red, right, empty circles = cold data). Orange oval (center) represents Maxwell demon performing promote/demote operations based on categorical distance.

Definition 5.6 (Categorical Neighborhood). The δ -neighborhood of a path π consists of all nodes reachable by modifying at most δ branch indices:

$$\mathcal{N}_\delta(\pi) = \{\pi' : d_H(\pi, \pi') \leq \delta\} \quad (39)$$

where d_H is the Hamming distance between paths.

Proposition 5.5 (Neighborhood Size). *The δ -neighborhood contains at most $\sum_{k=0}^{\delta} \binom{d}{k} 2^k$ paths, where d is the path length.*

Nearest neighbor search explores this neighborhood, collecting data from all nodes within categorical distance δ of the target.

5.7 Compression

The hierarchy can be compressed by removing unused branches.

Definition 5.7 (Hierarchy Compression). Compression removes all nodes that:

1. Hold no data, and
2. Have no descendants that hold data.

Proposition 5.6 (Compression Ratio). *If N_{data} nodes hold data out of N_{total} total nodes, the compression ratio is:*

$$R = \frac{N_{total}}{N_{compressed}} \geq \frac{N_{total}}{N_{data} \cdot D} \quad (40)$$

where D is the maximum depth and $N_{compressed}$ is the number of nodes after compression.

Compression preserves all stored data while eliminating structural overhead. For sparse storage patterns (few data items in a large potential space), compression can achieve high ratios.

5.8 Scale Ambiguity

The hierarchy exhibits scale ambiguity: the local structure is identical at every depth.

Theorem 5.7 (Scale Invariance). *For any depth d , the sub-hierarchy rooted at depth d is isomorphic to the full hierarchy rooted at depth 0:*

$$\mathcal{H}_d \cong \mathcal{H}_0 \quad (41)$$

where $\mathcal{H}_d$ denotes the hierarchy structure starting from any node at depth d .

Proof. Both structures are infinite 3-ary trees. The isomorphism is given by the identity map on the branching structure. Coordinate values differ by the depth-dependent scaling factor 3^{-d} , but this is a uniform rescaling that preserves all structural relationships. $\square$

Scale ambiguity means that an observer at depth d cannot determine their absolute depth from local measurements alone. The coordinate system looks the same at every scale.

6 Categorical Memory Architecture

6.1 Memory Architecture Overview

The categorical memory system integrates the S-entropy addressing, precision-by-difference navigation, hardware oscillation capture, and hierarchical storage into a complete memory architecture.

Definition 6.1 (Categorical Memory System). A categorical memory system $\mathcal{M}$ consists of:

1. An oscillator capture module $\mathcal{O}$ that samples hardware timing
2. A precision calculator $\mathcal{P}$ that computes precision-by-difference values
3. A categorical hierarchy $\mathcal{H}$ that provides the storage structure
4. A memory controller $\mathcal{C}$ that manages data placement and retrieval

6.2 Address Space

The address space of categorical memory is the S-entropy coordinate space.

Definition 6.2 (S-Entropy Address). An S-entropy address $\mathcal{A}$ consists of:

1. A precision-by-difference trajectory $\mathcal{T} = \{\Delta P(1), \dots, \Delta P(K)\}$
2. A trajectory hash $h = \mathcal{H}(\mathcal{T})$ providing a compact identifier
3. The current S-coordinate $\mathbf{S} = (S_k, S_t, S_e)$
4. The hierarchy path π derived from the trajectory

Unlike conventional addresses that are assigned externally, S-entropy addresses emerge from the access pattern itself. The trajectory of precision-by-difference values accumulated during access operations determines the address.

Proposition 6.1 (Address Uniqueness). *Two S-entropy addresses are equal if and only if their trajectories are identical:*

$$\mathcal{A}_1 = \mathcal{A}_2 \iff \mathcal{T}_1 = \mathcal{T}_2 \quad (42)$$

6.3 Storage Operations

6.3.1 Write Operation



Storing data at a categorical address proceeds as follows:

Require: Data D , key κ


- 1: Create new S-entropy address $\mathcal{A}$
- 2: Record initial precision sample: $\Delta P(1) \leftarrow \text{precision_sample}()$
- 3: Update address with sample: $\mathcal{A}.\text{record}(\Delta P(1))$
- 4: Repeat sampling to build trajectory of length K
- 5: Compute hierarchy path: $\pi \leftarrow \text{path_from_trajectory}(\mathcal{T})$
- 6: Navigate to node: $\nu \leftarrow \text{navigate}(\mathcal{H}, \pi, \text{create}=\text{true})$
- 7: Store data: $\nu.\text{data} \leftarrow D$
- 8: Store key: $\nu.\text{key} \leftarrow \kappa$
- 9: Index by key: $\text{key_index}[\kappa] \leftarrow \nu$
- 10: **return** $\mathcal{A}$

The trajectory is built dynamically during the write operation. The hardware timing during the write process determines the address.


Categorical Addressing: 3^k Hierarchy Structure S-Entropy Navigation and Coordinate Decomposition



C. Path Decomposition (Trajectory $\rightarrow$ Node Sequence)



D. Coordinate Decomposition (S-Space Partitioning)



(B) Node representation with S-coordinate ranges (bar chart, 12 nodes) shows unique coordinate assignment. Y-axis: data nodes (data_0 to data_11). X-axis: coordinate range [0, 1]. Each node displays three bars: S_k (blue, knowledge entropy), S_t (purple, temporal entropy), S_e (orange, evolution entropy). Depth labels (right): $d = 6$ to $d = 17$. Coordinate ranges non-overlapping, validating unique addressing. Example: data_0 ($d = 6$): $S_k \in [0.0, 0.2]$, $S_t \in [0.0, 0.2]$, $S_e \in [0.0, 0.2]$. Validates S-coordinate space provides complete addressing scheme.

(C) Path decomposition (trajectory $\rightarrow$ node sequence) shows address construction. Address: "alpha" (trajectory hash: 3b224a503f8397ec). 8 steps (0-7) with branch selection at each step. Step 0: Branch 0 (green), Path: [0], Region: 3^{-1} . Step 1: Branch 2 (red), Path: [02], Region: 3^{-2} . Step 2: Branch 2 (red), Path: [022], Region: 3^{-3} . Step 3: Branch 1 (orange), Path: [0221], Region: 3^{-4} . Step 4: Branch 0 (green), Path: [02210], Region: 3^{-5} . Step 5: Branch 2 (red), Path: [022102], Region: 3^{-6} . Step 6: Branch 2 (red), Path: [0221022], Region: 3^{-7} . Step 7: Branch 1 (orange), Path: [02210221], Region: 3^{-8} . Legend: ΔP branch selection (0 = $\Delta P > 0$, 1 = $\Delta P = 0$, 2 = $\Delta P < 0$). Gray arrows indicate sequential progression. Validates trajectory-based addressing where path history uniquely identifies location.

(D) Coordinate decomposition (S-space partitioning, 3D scatter) shows 30 points in (S_k, S_t, S_e) space. Axes: S_k (Knowledge, 0-1), S_t (Temporal, 0-1), S_e (Entropy, 0-1). Points colored by hierarchy depth (0.0-20.0 scale, blue to yellow gradient). Points cluster along trajectory path, forming curved structure in 3D space. Validates S-space partitioning where categorical distance (depth) corresponds to Euclidean distance in coordinate space.

(B) Node representation with S-coordinate ranges (bar chart, 12 nodes) shows unique coordinate assignment. Y-axis: data nodes (data_0 to data_11). X-axis: coordinate range [0, 1]. Each node displays three bars: S_k (blue, knowledge entropy), S_t (purple, temporal entropy), S_e (orange, evolution entropy). Depth labels (right): $d = 6$ to $d = 17$. Coordinate ranges

6.3.2 Read Operation

Retrieving data by S-entropy address:

Require: S-entropy address $\mathcal{A}$

- 1: Extract path: $\pi \leftarrow \mathcal{A}.\text{hierarchy_path}$
- 2: Navigate to node: $\nu \leftarrow \text{navigate}(\mathcal{H}, \pi, \text{create=false})$
- 3: **if** ν is null or $\nu.\text{data}$ is null **then**
- 4: **return** null (data not found)
- 5: **end if**
- 6: Update access count: $\nu.\text{access_count} \leftarrow \nu.\text{access_count} + 1$
- 7: **return** $\nu.\text{data}$

6.3.3 Read by Key

For applications requiring key-based access, a supplementary index provides $O(1)$ lookup:

Require: Key κ

- 1: $\nu \leftarrow \text{key_index}[\kappa]$
- 2: **if** ν is null **then**
- 3: **return** null
- 4: **end if**
- 5: **return** $\nu.\text{data}$

The key index provides conventional access semantics while the categorical structure provides semantic organization.

6.4 Memory Tiers

The categorical memory system supports multiple storage tiers with different performance characteristics. Figure 6(D) visualizes the exponential capacity growth across tiers, from L1 cache ($\sim 10^0$ items) to Archive ($\sim 10^9$ items).

Definition 6.3 (Memory Tier). A memory tier $\mathcal{T}_i$ is characterized by:

- Access latency ℓ_i
- Capacity C_i
- Energy cost per access E_i

Standard tiers in order of decreasing performance:

1. L1 Cache: $\ell \approx 1$ ns, $C \approx 64$ KB
2. L2 Cache: $\ell \approx 10$ ns, $C \approx 256$ KB
3. RAM: $\ell \approx 100$ ns, $C \approx$ GB
4. SSD: $\ell \approx 10^5$ ns, $C \approx$ TB
5. Archive: $\ell \approx 10^8$ ns, $C \approx$ unlimited

Figure 7.4.3(A) demonstrates the operational tier hierarchy, showing how categorical prefetching achieves 100% L1 hit rate by placing data in the correct tier before access.

Definition 6.4 (Tier Assignment). The tier of a datum is determined by its categorical position relative to the current access position:

$$\text{tier}(D) = f(d_S(\mathbf{S}_D, \mathbf{S}_{\text{current}})) \quad (43)$$

where d_S is the S-entropy distance and f is a monotonically increasing function mapping distance to tier index.

Data categorically close to current access patterns resides in fast tiers; data categorically distant resides in slow tiers.

6.5 Prediction and Prefetching

The trajectory completion mechanism enables predictive data movement.

Definition 6.5 (Access Prediction). Given the current trajectory $\mathcal{T}$, the predicted future accesses are data items whose S-coordinates are close to the trajectory completion point:

$$\text{predicted} = \{D : d_S(\mathbf{S}_D, \mathbf{S}^*) < \epsilon\} \quad (44)$$

where $\mathbf{S}^*$ is the predicted completion point.

```

1: function PREFETCH( $\mathcal{T}, n$ )
2:    $\mathbf{S}^* \leftarrow \text{predict\_completion}(\mathcal{T})$ 
3:   candidates  $\leftarrow \text{find\_nearest}(\mathbf{S}^*, n)$ 
4:   for each  $D$  in candidates do
5:     if tier( $D$ ) > 0 then                                ▷ Not already in fastest tier
6:       promote( $D$ , tier=0)
7:     end if
8:   end for
9: end function

```

Prefetching moves predicted data to fast tiers before it is actually accessed, reducing access latency.

6.6 Automatic Clustering

A key property of categorical memory is automatic semantic clustering.

Proposition 6.2 (Clustering Property). *Data accessed in similar patterns will have similar S-entropy addresses and thus reside in nearby positions in the hierarchy.*

Proof. Similar access patterns produce similar precision-by-difference trajectories (the timing environment during access is similar). Similar trajectories produce similar paths through the hierarchy (each precision value maps deterministically to a branch index). Similar paths lead to nearby nodes in the hierarchy. Therefore, data with similar access patterns clusters spatially in the hierarchy. $\square$

This clustering emerges without explicit indexing or classification. The hierarchical structure induced by precision-by-difference navigation naturally organizes data by access pattern similarity.

6.7 Statistics and Monitoring

The memory system maintains statistics for performance monitoring.

Definition 6.6 (Memory Statistics). Key statistics include:

- Tier occupancy: number of items in each tier
- Hit rates: fraction of accesses satisfied from each tier
- Navigation depth: average path length to accessed data
- Prediction accuracy: fraction of predicted accesses that occur
- Clustering quality: intra-cluster vs. inter-cluster distance ratio

These statistics enable tuning of tier capacities, prediction parameters, and prefetch aggressiveness.

7 Memory Controller as Maxwell Demon

7.1 Controller as Maxwell Demon

The categorical memory controller operates as a Maxwell demon in information space. The controller observes precision-by-difference values (categorical information) and uses them to sort data across memory tiers (hot vs. cold), analogous to the original demon sorting molecules by velocity (fast vs. slow). Figure 6(F) illustrates the demon schematically, showing how it mediates between fast tier (hot data, filled circles) and slow tier (cold data, empty circles).

Definition 7.1 (Categorical Memory Demon). The categorical memory demon $\mathcal{D}$ is an agent that:

1. Observes precision-by-difference values from hardware oscillators
2. Computes S-entropy coordinates from precision signatures
3. Determines categorical proximity of stored data to current position
4. Moves data between tiers based on proximity (promotion/demotion)
5. Predicts future access patterns through trajectory completion

The demon’s performance is validated in Figure 7.4.3(C), which shows 157 precision-by-difference calculations, 33 total hits, and zero evictions—demonstrating that categorical completion correctly predicts optimal tier placement.

7.2 Thermodynamic Considerations

The demon’s operation is consistent with thermodynamic laws because it operates in categorical space, which is orthogonal to physical phase space.

Theorem 7.1 (Demon Energy Cost). *The categorical memory demon incurs zero additional thermodynamic cost for categorical observation. The only energy costs are:*

1. *Physical tier transitions (moving bits requires energy)*
2. *Hardware timing measurements (negligible compared to normal operation)*

Proof. By Theorem 2.3, categorical observables commute with physical observables. Measuring S-entropy coordinates does not disturb physical phase space and therefore incurs no thermodynamic cost beyond the measurement apparatus itself.

The precision-by-difference values are derived from hardware timing that occurs regardless of the memory controller’s operation. The demon merely observes and interprets timing variations that are already present; it does not create them.

Energy is expended only when data physically moves between tiers (reading from one location, writing to another). This cost is identical to any tier management system. The demon’s decisions about *which* data to move incur no additional cost. $\square$

7.3 Controller State

The controller maintains state for navigation and management decisions.

Definition 7.2 (Controller State). The controller state $\mathcal{S}_c$ consists of:

- Current S-entropy position: $\mathbf{S}_{\text{current}}$
- Active S-entropy address: $\mathcal{A}_{\text{current}}$

- Tier contents: mapping from tier index to set of stored items
- Tier capacities: maximum items per tier
- Access statistics: hit/miss counts, latencies

The current position evolves with each memory operation as new precision-by-difference values are recorded.

7.4 Tier Management Policies

7.4.1 Promotion

Promotion moves data from a slower tier to a faster tier.

Definition 7.3 (Promotion Criterion). A datum D is promoted when its completion probability increases significantly:

$$P_{\text{completion}}(D, t) > \alpha \cdot P_{\text{completion}}(D, t - \Delta t) \quad (45)$$

where $\alpha > 1$ is the promotion threshold (typically $\alpha = 1.5$).

Completion probability measures how likely the datum is to be the categorical endpoint of the current trajectory. Increasing probability indicates the access pattern is converging toward this datum.

```

1: function CONSIDERPROMOTION(datum  $D$ )
2:    $P_{\text{new}} \leftarrow \text{completion\_probability}(D, \mathcal{A}_{\text{current}})$ 
3:   if  $P_{\text{new}} > \alpha \cdot D.\text{completion\_prob}$  then
4:      $\text{tier}_{\text{new}} \leftarrow D.\text{tier} - 1$  ▷ Move to faster tier
5:     if  $\text{tier}_{\text{new}} \geq 0$  then
6:        $\text{move}(D, \text{tier}_{\text{new}})$ 
7:     end if
8:   end if
9:    $D.\text{completion\_prob} \leftarrow P_{\text{new}}$ 
10: end function

```

7.4.2 Demotion

Demotion moves data from a faster tier to a slower tier, typically to make room for promoted data.

Definition 7.4 (Demotion Criterion). A datum D is selected for demotion when:

1. Its current tier is at capacity, and
2. It has the lowest completion probability among items in that tier.

```

1: function SELECTFORDEMOTION(tier  $\mathcal{T}$ )
2:    $P_{\text{min}} \leftarrow \infty$ 
3:   victim  $\leftarrow \text{null}$ 
4:   for each  $D$  in  $\mathcal{T}$  do
5:      $P \leftarrow \text{completion\_probability}(D, \mathcal{A}_{\text{current}})$ 
6:     if  $P < P_{\text{min}}$  then
7:        $P_{\text{min}} \leftarrow P$ 
8:       victim  $\leftarrow D$ 
9:     end if

```

```

10:   end for
11:   return victim
12: end function

```

This policy demotes data that is categorically unlikely to be accessed, based on the current trajectory, rather than simply the least recently used data.

7.4.3 Eviction

When the slowest tier reaches capacity, data must be evicted entirely.

Definition 7.5 (Eviction Policy). Eviction removes data from the system entirely. Selection criteria:

1. Already in the slowest tier (cannot be demoted further)
2. Lowest completion probability (least likely to be needed)
3. Oldest last access time (as tiebreaker)

7.5 Completion Probability Calculation

The completion probability quantifies how likely a datum is to be the endpoint of the current trajectory.

Definition 7.6 (Completion Probability). For a datum D with S-coordinate $\mathbf{S}_D$ and a trajectory with predicted completion $\mathbf{S}^*$:

$$P_{\text{completion}}(D) = \exp(-d_S(\mathbf{S}_D, \mathbf{S}^*)) \quad (46)$$

where d_S is the S-entropy distance.

This exponential decay assigns high probability to data close to the predicted completion point and low probability to distant data.

7.6 Categorical Distance as Tier Determinant

Proposition 7.2 (Distance-Tier Relationship). *The optimal tier for a datum D is:*

$$\text{tier}^*(D) = \min(T_{\max}, \lfloor c \cdot d_S(\mathbf{S}_D, \mathbf{S}_{\text{current}}) \rfloor) \quad (47)$$

where $T_{\max}$ is the slowest tier index and c is a scaling constant.

Data at zero categorical distance (current position) belongs in the fastest tier. Data at increasing distance belongs in progressively slower tiers.

7.7 Controller Algorithm


The complete controller algorithm integrates position tracking, tier management, and prediction:

```

1: function CONTROLLERSTEP
2:   Update current position from precision samples
3:   Compute completion predictions for active trajectories
4:   for each datum  $D$  accessed since last step do
5:     Update  $D.\text{access\_count}$ 
6:     ConsiderPromotion( $D$ )

```

Categorical Memory Operations: Maxwell Demon Controller Tier Management • Prefetching • Categorical Completion



(B) Storage & retrieval operations (bar chart, 20 hierarchy depths) shows uniform distribution. X-axis: hierarchy depth (6-20). Left y-axis: items stored (0-6). Right y-axis: normalized scale (0-1.0). Blue bars: stored items (1 per depth, uniform height). Green stars: retrieved items (4 retrievals at depths 6, 10, 16, 20). Annotation box (top-right): "Nodes: 258, Data: 20." Validates uniform storage distribution across hierarchy depths, with selective retrieval at specific depths.

(C) Maxwell demon controller performance (horizontal bar chart, 4 metrics) shows zero-eviction operation. Evictions: 0 (no bar). Promotions: 0 (no bar). Total Hits: 33 (green bar, short). Active Adrrs: 31 (purple bar, short). Total Calcs: 157 (blue bar, longest, annotation: "Mean ΔP : 2.81×10^{-6} s, Std ΔP : 4.90×10^{-7} s"). Zero evictions and promotions validate categorical completion predicts optimal tier placement, eliminating traditional cache management overhead. Mean ΔP of 2.81 μ s indicates sub-microsecond precision for trajectory calculations.

(D) Branch usage & navigation paths (stacked bar chart, 3 levels) shows balanced branching. X-axis: hierarchy level (0, 1, 2). Y-axis: branch count (0-12). Branch 0 (green), Branch 1 (orange), Branch 2 (red). Level 0: 3 total (1+1+1). Level 1: 8 total (2+5+1). Level 2: 13 total (4+8+1). Annotation box (top): "Navigation Paths: alpha→beta: 40 steps, alpha→gamma: 38 steps, beta→gamma: 40 steps." Branch 1 (orange, $\Delta P = 0$) most frequent at all levels, indicating temporal synchronization events dominate navigation. Validates balanced branch usage where all three ΔP conditions (> 0 , $= 0$, < 0) contribute to addressing.

(B) Storage & retrieval operations (bar chart, 20 hierarchy depths) shows uniform distribution. X-axis: hierarchy depth (6-20). Left y-axis: items stored (0-6). Right y-axis: normalized scale (0-1.0). Blue bars: stored items (1 per depth, uniform height). Green stars: retrieved items (4 retrievals at depths 6, 10, 16, 20). Annotation box (top-right): "Nodes: 258, Data: 20." Validates uniform storage distribution across hierarchy depths, with selective retrieval at specific depths.

```

7:   end for
8:   if any tier at capacity then
9:       Demote lowest-probability items
10:  end if
11:  Prefetch high-probability items not in fast tiers
12: end function

```

The controller runs continuously, updating the memory organization in response to evolving access patterns.

7.8 Synchronization

Multiple concurrent accessors require synchronization to maintain consistent controller state.

Definition 7.7 (Controller Lock). A reentrant lock protects controller state modifications. All operations that modify tier contents, position, or statistics acquire this lock.

The lock granularity is chosen to balance consistency with concurrency. Fine-grained locking per tier enables parallel access to different tiers; coarse-grained locking on the entire controller simplifies reasoning about state consistency.

8 Experimental Validation

The categorical memory architecture has been validated through comprehensive experimental implementation. The validation framework measures precision-by-difference values from real hardware oscillators, constructs S-entropy coordinates, and demonstrates memory operations on the resulting hierarchical structure.

8.1 Hardware Oscillation Capture

Real hardware timing measurements were captured from multiple oscillator sources: CPU performance counters, memory access latency, and computation jitter. Figure 4.6 demonstrates the complete measurement pipeline, showing how hardware oscillations at frequencies ranging from 50 Hz (power supply) to 8 GHz (PCIe bus) are sampled and converted to precision-by-difference values.

The captured timing variations exhibit a mean precision-by-difference of $\bar{\Delta P} = 0.0086$ ms with standard deviation $\sigma_{\Delta P} = 0.193$ ms, providing sufficient dynamic range for S-entropy coordinate encoding. As shown in Figure 4.6(B), the oscillation harvesting process produces stable ΔP values across multiple hardware sources.

8.2 Precision-by-Difference Network

The precision-by-difference network achieves 96.1% latency reduction compared to traditional request-response communication, as demonstrated in Figure 3.3. The temporal coherence windows maintain 99.90% quality, validating the synchronization mechanism underlying S-entropy navigation.

Key validation results include:

- Branch selection: balanced ternary distribution (31.3% / 38.1% / 30.6%)
- Hierarchy coverage: 22.0% of $3^5 = 243$ nodes accessed during navigation
- Prediction accuracy: mean error 0.0078 for trajectory completion

8.3 Memory Operations

The categorical memory controller achieves 100% hit rate with zero evictions, as shown in Figure 7.4.3. All 33 accesses were satisfied by L1 cache, validating that categorical prefetching correctly places data in the optimal tier before access occurs.

The Maxwell demon controller computed 157 precision-by-difference calculations with mean $\Delta P = 2.81 \times 10^{-6}$ s and standard deviation 4.90×10^{-7} s, demonstrating sub-microsecond precision for trajectory calculations.

9 Discussion

9.1 Relationship to Maxwell’s Demon

The memory controller described in this work operates as a Maxwell demon in a precise sense. Maxwell’s original thought experiment [5] proposed an intelligent agent that could sort molecules by velocity, apparently decreasing entropy without work. Modern resolution of this paradox, through the work of Landauer [3] and Bennett [1], established that the demon must expend energy to erase information, preserving the second law.

Our memory controller demon operates in information space rather than physical space. It observes precision-by-difference values (categorical observables) rather than molecular velocities (physical observables). As illustrated in Figure 2.1, the categorical and physical faces of information are complementary—Maxwell observed only the kinetic face, missing the categorical structure entirely. The “demon” is revealed as a projection artifact: the shadow of categorical dynamics onto the observable kinetic plane.

The key insight from our categorical measurement framework is that categorical observables commute with physical observables:

$$[\hat{O}_{\text{categorical}}, \hat{O}_{\text{physical}}] = 0 \quad (48)$$

This commutation relation implies that the demon can extract information about categorical position without disturbing physical state. The information gain comes not from physical measurement but from observing the inherent timing structure of hardware oscillations.

The demon’s operation—promoting data to faster tiers when categorical completion probability increases, demoting when it decreases—does not violate thermodynamics because it operates on categorical coordinates orthogonal to physical phase space. The energy cost of memory tier transitions is physical (moving bits requires energy), but the decision of which data to move is categorical (based on S-entropy position) and incurs no additional thermodynamic cost. This is validated experimentally in Figure 7.4.3(C), which shows zero evictions and zero promotions when the categorical completion mechanism correctly predicts optimal tier placement.

9.2 Scale Ambiguity and Recursive Self-Similarity

A notable property of the 3^k hierarchical structure is scale ambiguity. The mathematical structure of S-entropy coordinates is identical at every hierarchical level. An observer at depth k cannot determine their absolute position in the hierarchy from local measurements alone—the coordinate system at depth k looks identical to the coordinate system at depth $k + 10$ or depth $k - 5$.

This property emerges from the recursive self-similarity of the branching structure. Each node, regardless of depth, branches into three children corresponding to the same three S-entropy dimensions. The transformation from parent to child coordinates follows the same rules at every level:

$$\mathbf{S}_{\text{child}}^{(i)} = \frac{1}{3}\mathbf{S}_{\text{parent}} + \epsilon_i, \quad i \in \{0, 1, 2\} \quad (49)$$

where ϵ_i is a perturbation in the i -th dimension.

This scale ambiguity is not a limitation but a feature. It means the addressing scheme is intrinsically relative rather than absolute. A datum’s address specifies its position relative to other data in the same computational context, not its position in some absolute coordinate frame. This relativity is appropriate for semantic organization, where meaning is contextual rather than absolute.

9.3 Comparison with Conventional Memory Hierarchies

Conventional memory hierarchies (L1 cache, L2 cache, RAM, SSD, etc.) use physical proximity and access recency to determine data placement. The LRU (Least Recently Used) policy and its variants make placement decisions based on when data was last accessed, with the assumption that recently accessed data is likely to be accessed again soon.

Categorical memory makes placement decisions based on categorical proximity rather than temporal recency. Data that is categorically close to the current position (small S-entropy distance) is placed in fast tiers; data that is categorically distant is placed in slow tiers. This approach has different characteristics:

1. **Pattern sensitivity:** Categorical placement is sensitive to access patterns, not just timing. Two data items accessed at the same time but in different patterns will have different categorical positions.
2. **Prediction vs. reaction:** Conventional caches react to past accesses. Categorical memory predicts future accesses based on trajectory completion, enabling proactive rather than reactive placement.
3. **Semantic clustering:** Data accessed in similar patterns cluster together in the hierarchy, even if they are physically unrelated. This provides implicit indexing based on usage semantics.

The $O(\log n)$ navigation complexity is higher than conventional $O(1)$ addressing, but the navigation process itself contributes information about data relationships that conventional addressing discards.

10 Conclusion

We have presented a memory architecture based on categorical addressing through S-entropy coordinates and precision-by-difference navigation. The central contributions of this work are:

1. **S-Entropy Address Space:** We established that the three-dimensional coordinate system $\mathbf{S} = (S_k, S_t, S_e)$ provides a complete addressing scheme for hierarchical memory, as illustrated in Figure 6.3.1. Each coordinate encodes a distinct aspect of categorical position: knowledge entropy (state uncertainty), temporal entropy (timing uncertainty), and evolution entropy (trajectory uncertainty). The address space has recursive 3^k structure with scale-ambiguous self-similarity, visualized in Figure 6.3.1(A).
2. **Precision-by-Difference as Address:** We demonstrated that the precision-by-difference value $\Delta P = T_{\text{ref}} - t_{\text{local}}$ encodes position in the categorical hierarchy. The accumulated sequence of precision-by-difference values forms a trajectory whose hash uniquely identifies a storage location. Figure 3.3 validates this mechanism, showing 96.1% latency reduction through trajectory-based navigation. The access history *is* the address.

3. **Hardware Grounding:** We showed that precision-by-difference values can be obtained from real hardware oscillations: CPU cycle timing, memory access latency, and I/O jitter. Figure 4.6 demonstrates the complete pipeline from hardware oscillation capture through S-entropy coordinate computation. These physical timing variations provide the raw material from which categorical coordinates are computed.
4. **Maxwell Demon Controller:** We described a memory controller that operates as a categorical Maxwell demon, navigating the S-entropy hierarchy to determine optimal data placement. Figure 7.4.3 demonstrates 100% hit rate with zero evictions, validating that the controller uses categorical completion to predict access patterns and achieve proactive tier assignment based on trajectory endpoints rather than reactive caching based on access recency.
5. **Experimental Validation:** We validated the architecture through comprehensive implementation, demonstrating $O(\log n)$ navigation complexity, automatic clustering of related data, and effective tier management through categorical completion. The validation results are summarized in Section 8 and visualized across Figures 6–7.4.3.

The categorical memory architecture provides a theoretical foundation for memory systems where addressing reflects meaning rather than position. The precision-by-difference mechanism grounds abstract categorical coordinates in physical hardware timing, while the Maxwell demon controller provides intelligent data placement based on categorical proximity rather than temporal recency. The experimental validation confirms that this architecture achieves its theoretical promises: zero-backaction categorical measurement, trajectory-based addressing, and thermodynamically consistent demon operation.

References

- [1] Charles H Bennett. The thermodynamics of computation—a review. *International Journal of Theoretical Physics*, 21(12):905–940, 1982.
- [2] John L Hennessy and David A Patterson. *Computer Architecture: A Quantitative Approach*. Morgan Kaufmann, 6th edition, 2017.
- [3] Rolf Landauer. Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3):183–191, 1961.
- [4] Harvey S Leff and Andrew F Rex. *Maxwell’s Demon 2: Entropy, Classical and Quantum Information, Computing*. CRC Press, 2002.
- [5] James Clerk Maxwell. Theory of heat. 1871.